

COMP 3500 Introduction to Operating Systems

Project 1 – Using the Linux Terminal

Version 3.6

Revised: January 13, 2023

Points Possible: 100

There should be no collaboration among students. A student shouldn't share any project code with any other student. Collaborations among students in any form will be treated as a serious violation of the University's academic integrity code.

Objectives:

- Get to know the Linux operating system
- Learn basic Linux commands
- Prepare a Linux programming environment for the future COMP3500 projects
- Learn how to remotely access Linux servers (see Option 1 below).
- Learn how to install VirtualBox and CentOS/Ubuntu (see Option 2 below).
- Learn how to install Docker and Ubuntu (see Option 3 below).
- Compile and debug your first C program in Linux

Requirements:

- Each student should **independently** accomplish this project assignment. You may discuss with other COMP3500 student to solve the coding problems.
- To embark on this project, you may choose one of the following four options.
 - Option 1: use SSH to connect to a remote Linux server. Please read "*Project 1 Option 1 Remotely Connect to Linux Machines.pdf*" for details.
 - Option 2: install VirtualBox first on your computer followed by the installation of Linux. Please read "*Project 1 Option 2 Install VirtualBox and CentOS.pdf*" for details. **Note:** You are allowed to install Ubuntu.
 - Option 3: Install Ubuntu Linux in a Docker Container. Please read "*Project 1 Option 3 Install Ubuntu Linux in a Docker container.pptx*" for details.
 - Option 4: Dual-boot Linux and Windows on your own laptop/desktop computers.
- You are highly recommended to use CentOS or Ubuntu as a Linux distribution.

1. Setup Linux Programming Environment (10 Points)

CAUTION! When you access a remote Linux machine in the Linux lab (i.e., Option 1), please *don't use the "sudo" command*, because you have no administrative rights. If you choose options 2-4, you may use the "sudo" command on your virtual machine or local machine where you are a system admin.

You should choose one of the following four options to setup your Linux programming environment.

1.1 Option 1: Remotely Connect to Linux Machines (Difficulty Level: ★)

Please follow the instructions specified in “*Project 1 Option 1 Remotely Connect to Linux Machines.pdf*” to learn how to access a remote Linux server. You must keep in mind that when you access a remote Linux machine in the Linux lab, please *don't use the "sudo" command*.

1.2 Option 2: Install VirtualBox and CentOS (Difficulty Level: ★★)

Note: You are allowed to install Ubuntu.

Please follow the instructions specified in “*Project 1 Option 2 Install VirtualBox and CentOS.pdf*” to learn a way of installing VirtualBox and CentOS (or Ubuntu) in your own laptop or desktop computers.

1.3 Option 3: Install Ubuntu Linux in a Docker container (Difficulty Level: ★★)

Note: You may install CentOS in the Docker container.

- Please download and follow the instructions below to kickoff this option. The instructions are available on Canvas.
Project 1 Option 3 Install Ubuntu Linux in a Docker container.pptx
- After you install the Ubuntu Linux in the Docker container, please follow the steps to finish the configuration.
- Open a terminal and do “yum -y update” to update all the packages. You need to be the root to do this.
- Please make sure the following packages are installed: **gcc, gcc-c++, vim-enhanced, emacs gdb, ethtool, hdparm, file** and **pciutils**. They most likely are already installed. You can run the following commands to confirm their availability and install as needed.

```
$yum -y install gcc
$yum -y install gcc-c++
$yum -y install vim-enhanced
$yum -y install emacs
$yum -y install gdb
$yum -y install ethtool
$yum -y install pciutils
$yum -y install file
```

1.4 Option 4: Dual-boot Linux and Windows (Difficulty Level: ★★★★★)

If you choose this option to dual boot Linux and Windows 10 on your own computer, please follow the instructions here:

<https://itsfoss.com/guide-install-linux-mint-16-dual-boot-windows/>

2. The `script` Command

The `script` command line tool allows you to save a session of your terminal. In addition to saving each command per line in a text file, the `script` command makes a typescript of everything that happens on your Linux terminal. Screencasting tools to a desktop session(GUI) is what `script` is to a terminal. Let us demonstrate the usage of `script` through the following example:

```
$ script
$Script started, file is typescript
$ cd
$ ls
file1 file2 file3
$ exit
exit
Script done, file is typescript
```

Then, you may use the `mv` command to change the file name from `typescript` to any name you like. Alternatively, you may specify the name of your log file upfront as below:

```
$ script sample.script
$Script started, file is typescript
$ cd
$ ls
file1 file2 file3
$ exit
exit
Script done, file is sample.script
```

3. Tasks (90 Points)

Script the following session using the `script` command. You may save each session (i.e., each task below) in one script file. Using the `tar` command to submit a tarred and compressed file named `project1.tgz` (see Section 4.2 for details).

3.1 (Task 1: 25 points. Difficulty Level: ★) Please use the `script` command to create a file named `"commands.script"` demonstrating that you understand how to use the following basic Linux commands.

- `man [command]` – displays the help information for the specified command.
- `cd` – changes a directory.
- `pwd` – displays the pathname for the current directory.
- `ls` – lists the files in the current directory.
- `mkdir` – makes a new directory.
- `cp` – copies a file from one location to another.
- `mv` – moves a file from one location to another.
- `rm` – removes a file.

- `rmdir [options] directory` – deletes empty directories.
- `chmod [options] mode filename` – changes a file's permissions.
- `clear` – clears a command line screen/window for a fresh start.
- `top` – displays the resources being used on your system. Press `q` to exit.
- `who [options]` – displays who is logged on.
- `nproc` – displays the number of cores.

Reference: How to Start Using the Linux Terminal

<https://www.howtogeek.com/140679/beginner-geek-how-to-start-using-the-linux-terminal/>

3.2 (Task 2: 10 points. Difficulty Level: ★) When you log in a Linux system, you should get to know the system in more detail. You are asked to find out your computer system's attributes, including **CPU frequency, cache size, memory size, the list of PCI devices, hard drive, network MAC address and link speed, and the devices generating interrupts**. The following system commands can help you. Many of these commands are for system administration, so you should run them as root if needed.

```
$more /proc/cpuinfo
$more /proc/meminfo
$more /proc/interrupts
$lspci
```

Tip: You may pipe the output of any Linux command to a file on Linux. For example:

```
#ls > test.txt
```

The above command writes all files in the current directory to a file called "test.txt" instead of displaying the files on a monitor. If you choose this option, you should include the pipeline output file(s) in your submitted tarball.

Please store the output from the above four commands into the following four files using the pipe (see the above tip).

- `cpuinfo.txt`
- `meminfo.txt`
- `lspci.txt`
- `interrupts.txt`

3.3 (Task 3: 30 points. Difficulty Level: ★★) With the computer up and running, you should give it a try to see if you can use the utilities on the system. For a system programmer, these include at least the editor, the compiler, the libraries, and the debugger. You are asked to do the following

Note: You must make use of the `script` command (see Section 2 page 3) to save a session for task 3. The script file name should be `"gdb.script"`.

3.3.1. Using your favorite editor, code a program (simple.c) that processes an array of 10 numbers, calculates the average of their square roots, and prints it out. I recommend either **emacs** or **vim** (a.k.a., vi) as an editor for COMP 3500. If you are not yet proficient with either of them, you should be able to type the `simple.c` program without too much effort.

3.3.2. The GNU compile is the default open source compiler on Linux. You should check a little on what gcc you have, and then compile your program as follows. Please do not forget the flag '-g' for using debugger in Step 4.

```
$gcc -v
$gcc -g -lm -o simple simple.c
```

3.3.3 Practice the command **ldd**, and understand the libraries on which your program is dependent upon for execution.

```
$man ldd
$ldd simple
```

3.3.4 The debugger is a friend you must get acquainted with to be a good programmer. Here is a little trick in using the GNU `gdb` debugger.

(1) First run your program, **simple**, alone

(2) Using Vim to create a file name as `.gdbinit` in the current directory. Your `.gdbinit` file should contains the following lines:

```
file simple
break main
break sqrt
info registers
```

(3) Run the `gdb` debugger and then test and try these commands ('r', 's', 'n', and 'c'), one command at a time complete the program.

```
# gdb
# r
# n
# s
# c
```

Note: You don't have to follow the above command sequence in `gdb`: you may use any order of your choice. You will receive full credit if you demonstrate that you are familiar with the basic `gdb` commands such as run, next, step, continue.

3.3.5 Your project report should include the source program, and the output from Steps (1)-(4).

3.4 (Task 4: 25 points. Difficulty Level: ★★) Now you should learn how to use `git` – a version control system - to manage your software development conducted by your future groups. An online book focusing on `git` can be found here: <https://git-scm.com/book/en/v2>.

Note: You must make use of the script command (see Section 2 page 3) to save a session for task 4. The script file name should be “`git.script`”.

3.4.1 Install git

CAUTION! This step is required for Options 2, 3, and 4 only. If you choose option 1, please skip this step and directly go to 3.4.2.

```
$ sudo yum install git
```

Reference: <https://www.linode.com/docs/development/version-control/how-to-install-git-on-linux-mac-and-windows/>

CAUTION! Please refer to the following webpage for detailed instructions on Steps 3.4.2-3.4.9.

<https://git-scm.com/book/en/v2/Git-Basics-Recording-Changes-to-the-Repository>

3.4.2 Initializing a Repository in an Existing Directory. We assume that you have created a directory “`comp3500/project1`” in your home directory.

```
$ cd ~/comp3500/project1
```

If you want to start version-controlling existing files (as opposed to an empty directory), you should probably begin tracking those files and do an initial commit. You can accomplish that with a few `git add` commands that specify the files you want to track, followed by a `git commit`:

```
$ git init
$ touch cpu.c
$ touch driver.c
$ touch file.c
$ git status
$ git add *.c
$ git status
$ git commit -m 'Initial project 1'
$ git status
$ rm cpu.c
$ ls
$ git checkout cpu.c
$ git checkout -b master2
$ touch cpu2.c
$ git add cpu2.c
$ git commit -m 'cpu2 is added into master 2'
```

```
$git branch
$ls
$git checkout master
$ls
$git clone -b master . ../new_project1
$cd ..
$ls
$cd new_project1
```

3.4.3 Cloning an Existing Repository

You clone a repository with `git clone <url>`. For example, if you want to clone the Git linkable library called `libgit2`, you can do so like this:

```
$git clone https://github.com/libgit2/libgit2
```

That creates a directory named `libgit2`, initializes a `.git` directory inside it, pulls down all the data for that repository, and checks out a working copy of the latest version. If you go into the new `libgit2` directory that was just created, you'll see the project files in there, ready to be worked on or used. If you want to clone the repository into a directory named something other than `libgit2`, you can specify the new directory name as an additional argument:

```
$git clone https://github.com/libgit2/libgit2 mylibgit
```

3.4.4 Checking the Status of Your Files:

```
$git status
```

3.4.5 Tracking New Files

In order to begin tracking a new file, you use the command `git add`. To begin tracking the `README` file, you can run this:

```
$git add README
```

3.4.6 Staging Modified Files

Now we change a file that was already tracked. If you change a previously tracked file called `CONTRIBUTING.md` and then run your `git status` command again.

3.4.7 Viewing Your Staged and Unstaged Changes

If the `git status` command is too vague for you — you want to know exactly what you changed, not just which files were changed — you can use the `git diff` command.

```
$git diff
```

3.4.8 Committing Your Changes

Now that your staging area is set up the way you want it, you can commit your changes. Remember that anything that is still unstaged — any files you have created or modified that you haven't run `git add` on since you edited them — won't go into this commit. They will stay as modified files on your disk. In this case, let's say that the last time you ran `git status`, you saw that everything was staged, so you're ready to commit your changes. The simplest way to commit is to type `git commit`:

```
$git commit
```

3.4.9 Viewing the Commit History

The most basic and powerful tool to do this is the `git log` command.

```
$git log
```

4. Deliverables

One of the following deliverables is acceptable. You may choose to submit a single PDF file (see Section 4.1) or a tarred and compressed file (see Section 4.2).

4.1 Project Report

Please submit your project report with needed contents as specified in the questions (see Section 3). You must submit your report through Canvas (no e-mail submission is accepted). Please report any problems you have solved when you learn Linux commands and `git`.

The file name should be formatted as:

```
"project1.pdf" or "project1.txt"
```

4.2 Multiple Script Files

If you have generated multiple script files, please save all the script files in one directory say (project1). Then, you should archive all the script files (see Sections 3.1, 3.3, 3.4) and text files (see Section 3.2) along with your report (See Section 4.1) and the source code (simple.c) into a single tarred and compressed file. Assume that the script files, text files, and your report are located in `~/comp3500/project1`, then you can follow the instructions below to prepare a single compressed file.

```
%tar vfcz project1.tgz ~/comp3500/project1
```

Your tarball should embrace the following files:

- project1.pdf or project1.txt (see Section 4.1)
- commands.script (see Section 3.1)

- `cpuinfo.txt` (see Section 3.2)
- `meminfo.txt` (see Section 3.2)
- `lspci.txt` (see Section 3.2)
- `interrupts.txt` (see Section 3.2)
- `gdb.script` (see Section 3.3)
- `git.script` (see Section 3.4)

5. Grading Criteria

- 1) Setting up your Linux programming environment: 10% (see Section 1)
- 2) Using Linux commands: 25% (see Section 3.1)
- 3) Getting to know your system: 10% (see Section 3.2)
- 4) Using `gcc` and `gdb`: 30% (see Section 3.3)
- 5) Using `git`: 25% (see Section 3.4)

6. Late Submission Penalty

- Ten percent (10%) penalty per day for late submission. For example, an assignment submitted after the deadline but up to 1 day (24 hours) late can achieve a maximum of 90% of points allocated for the assignment. An assignment submitted after the deadline but up to 2 days (48 hours) late can achieve a maximum of 80% of points allocated for the assignment.
- Assignment submitted more than 3 days (72 hours) after the deadline will not be graded.

7. Rebuttal period

- You will be given a period of **one week** to read and respond to the comments and grades of your homework or project assignment. The TA may use this opportunity to address any concern and question you have. The TA also may ask for additional information from you regarding your homework or project.